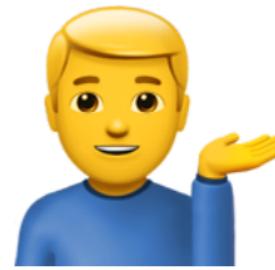


Docker: The Efficient

Introduction



Who am I?

- ▶ My name is Jannis Seemann
- ▶ I'm German , with a degree in computer science (TU München)
- ▶ **Important work experience:**
 - ▶ With 18, first internship at Google in London 
 - ▶ At age 20, another one in Mountain View, California 
 - ▶ Nowadays, I'm a freelance web developer (and server admin)
 - ▶ And one of the most popular instructors on the German speaking market on Udemy (300.000+ students)
- ▶ **In this course:**
 - ▶ You benefit from my work experience
 - ▶ ... and from my online teaching experience



Day 1

Understand container lifecycles

Apply essential CLI commands for containers and images

Build custom images using **Dockerfiles**



Day 2

Manage storage with volumes and bind mounts

Connect and isolate containers through custom networks

Orchestrate multi-container apps with Docker Compose



Day 2.5

Set up and manage private docker registries

Automate everything: From GitHub to automatic deployment

Create a fully-fledged deployment- & dev setup

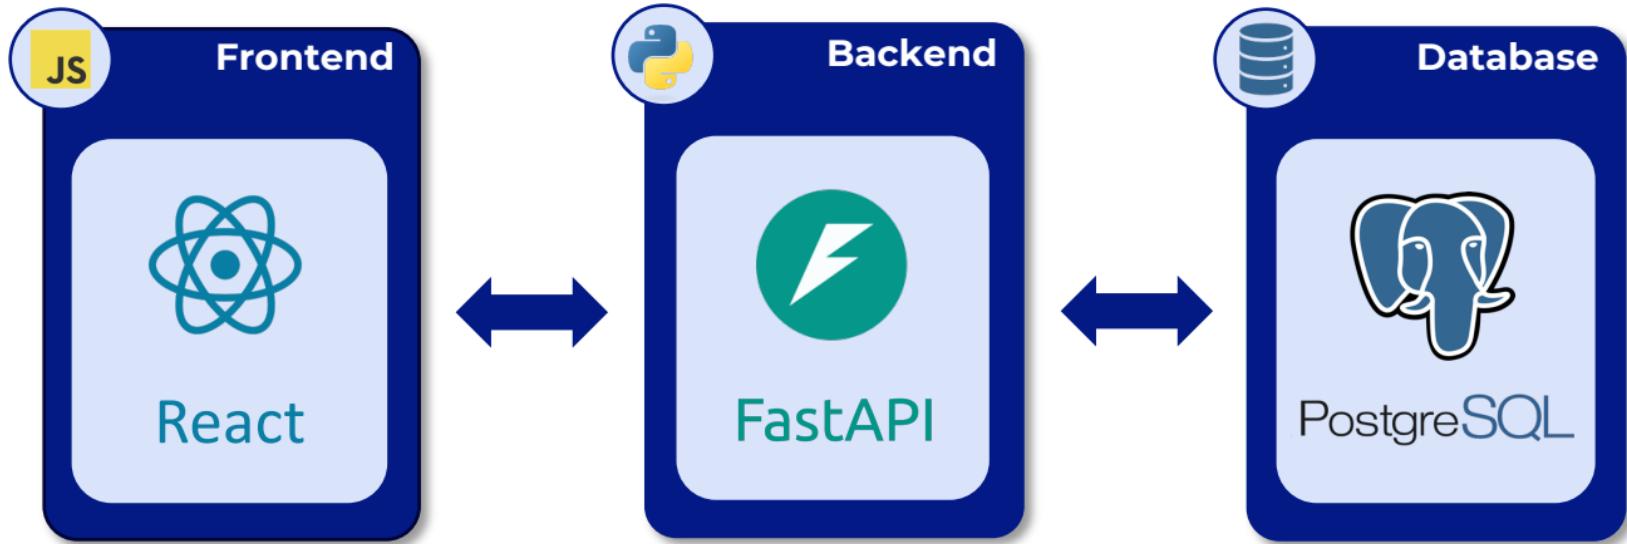
Thank you

- ▶ Thank you for subscribing / purchasing this course
- ▶ This really means a lot to us...
- ▶ ... as this allows us to dedicate so much time to our courses

Docker: The Efficient

See what Docker does!





Which pain points does Docker solve?

 Simplified setup & portability	 Cross-platform consistency	 Scalability & reproducibility
No manual installations or compatibility issues	Works seamlessly on Windows, macOS, and Linux	Easily scale services by simple configuration changes
Share and run the app anywhere with just a <code>compose.yml</code> file without extra setup	Guarantees consistent dependencies and avoids version mismatches	Ensures identical environments across development, testing, and production



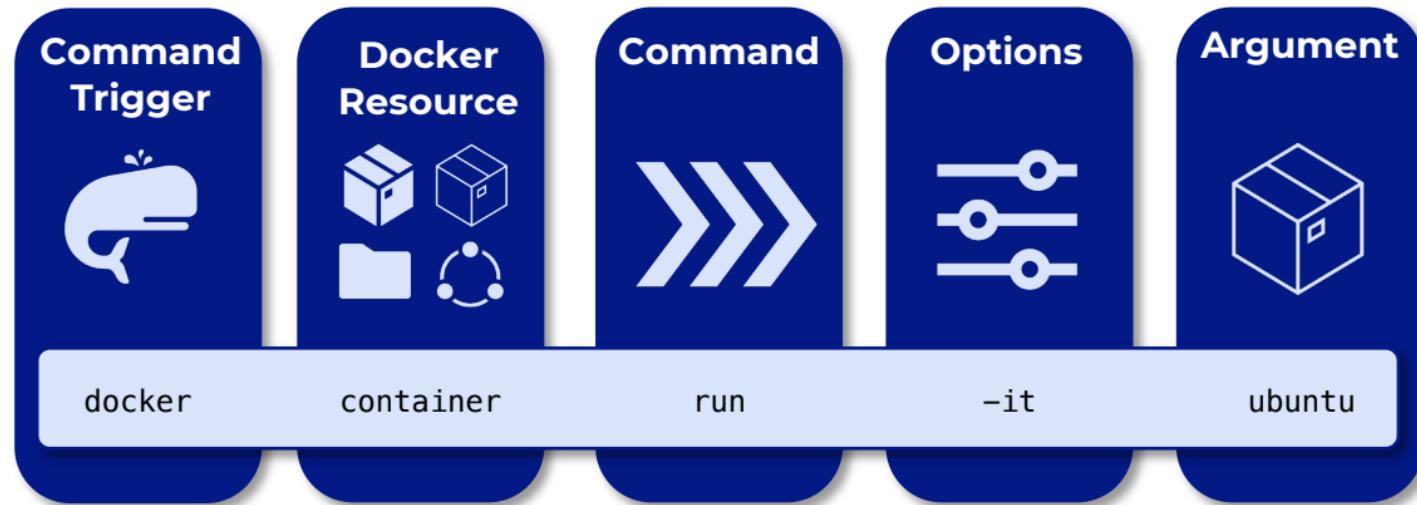
Docker replaces manual setups with a single command that works consistently everywhere!

Docker: The Efficient Guide

Run your first Docker containers!



The structure of a Docker command





Challenge

Run a `node` container and verify it by executing the following commands:

```
console.log(0.1+0.2==0.3);
```

```
console.log([...Array(5).keys()]);
```

```
console.log(typeof NaN);
```

Docker: The Efficient

What is a container?



What is a Container?



Efficient runtime environment

Provides everything an application needs to run

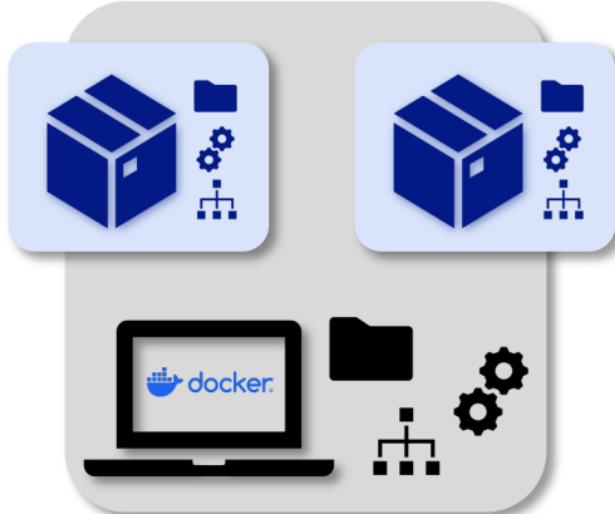
Uses fewer resources than virtual machines



Isolated and self-sufficient

Isolated from the rest of your system and other containers

Runs independently, with its own filesystem, processes, and network stack



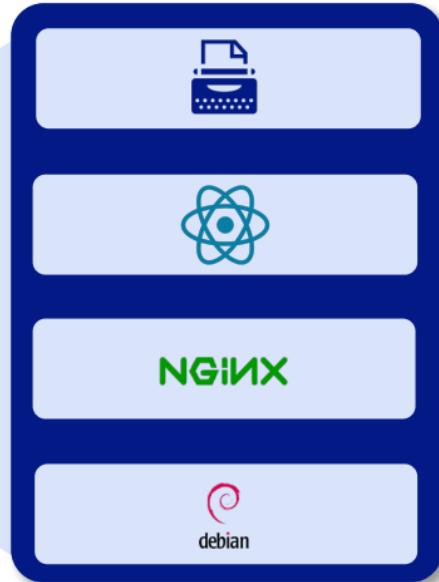
The structure of a container



A container packages the entire code base and all dependencies of an application



Enables it to run consistently across different systems

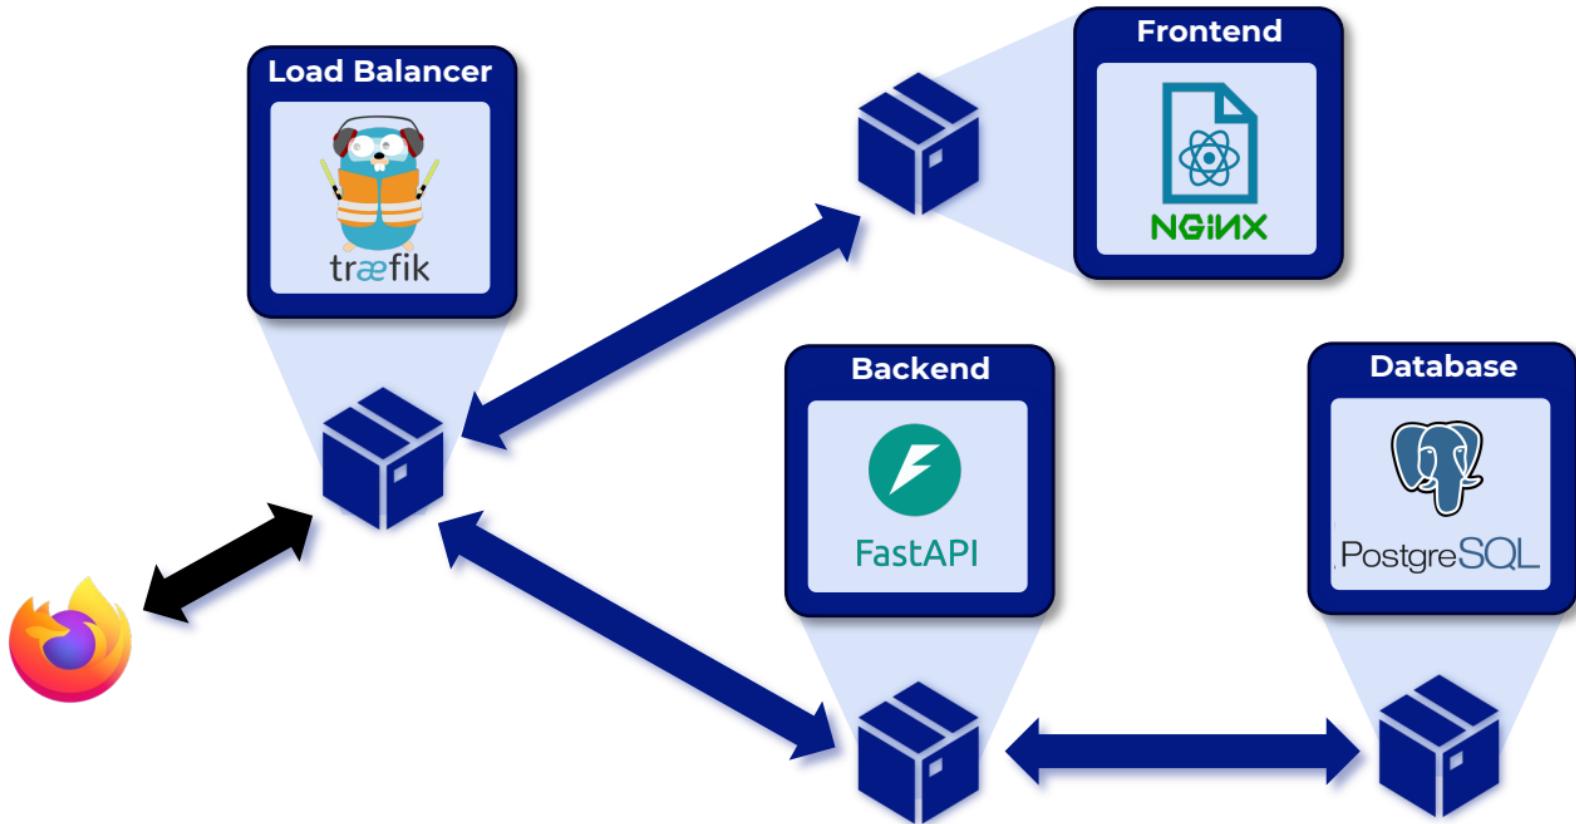


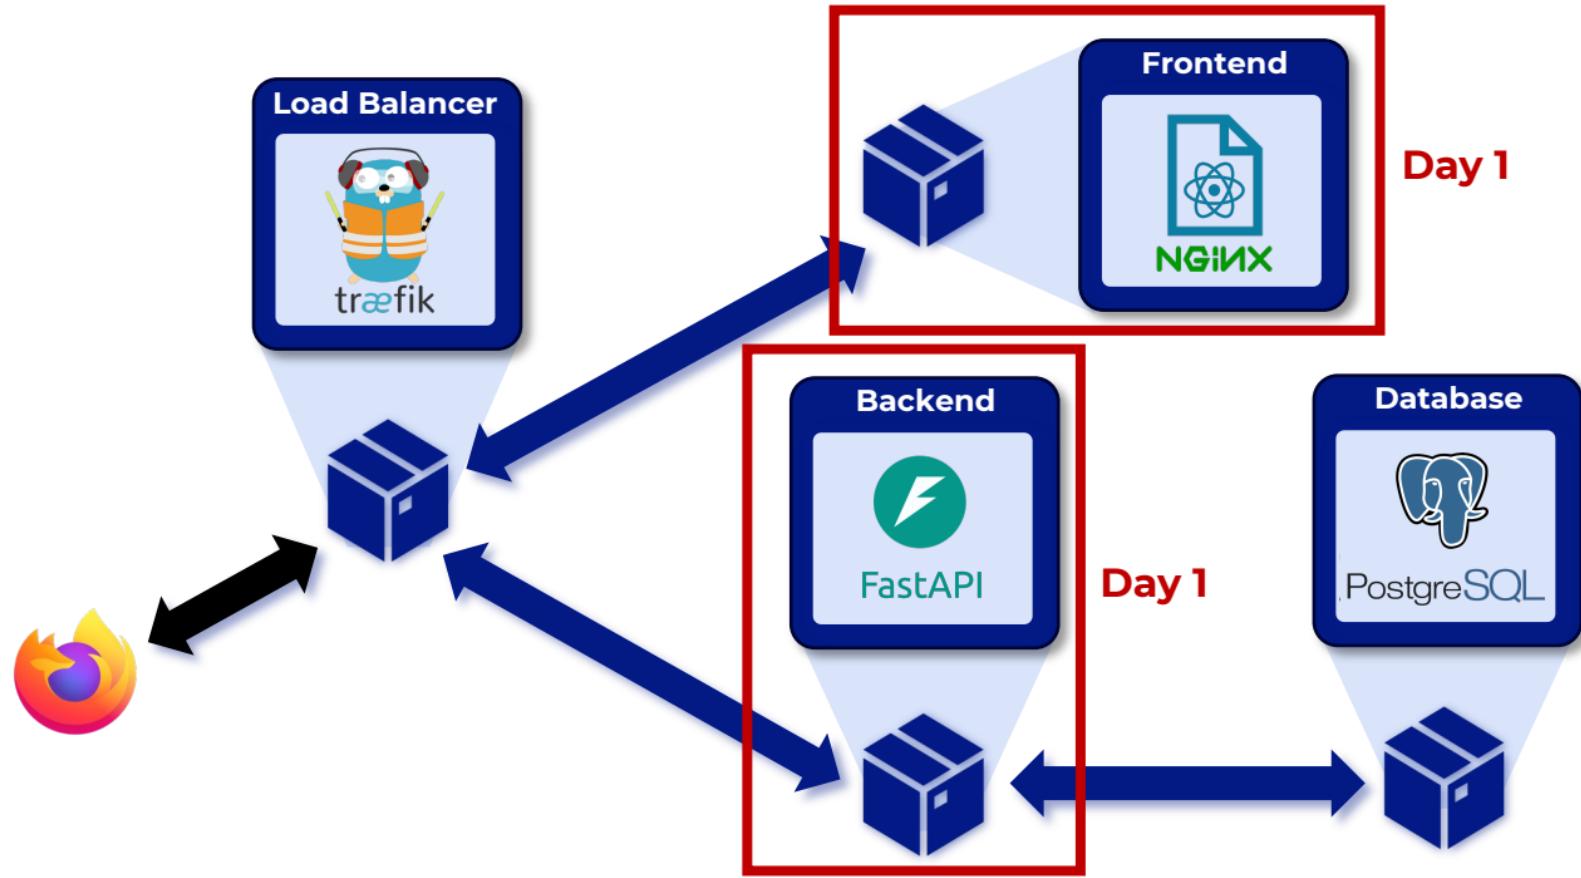
Writable layer

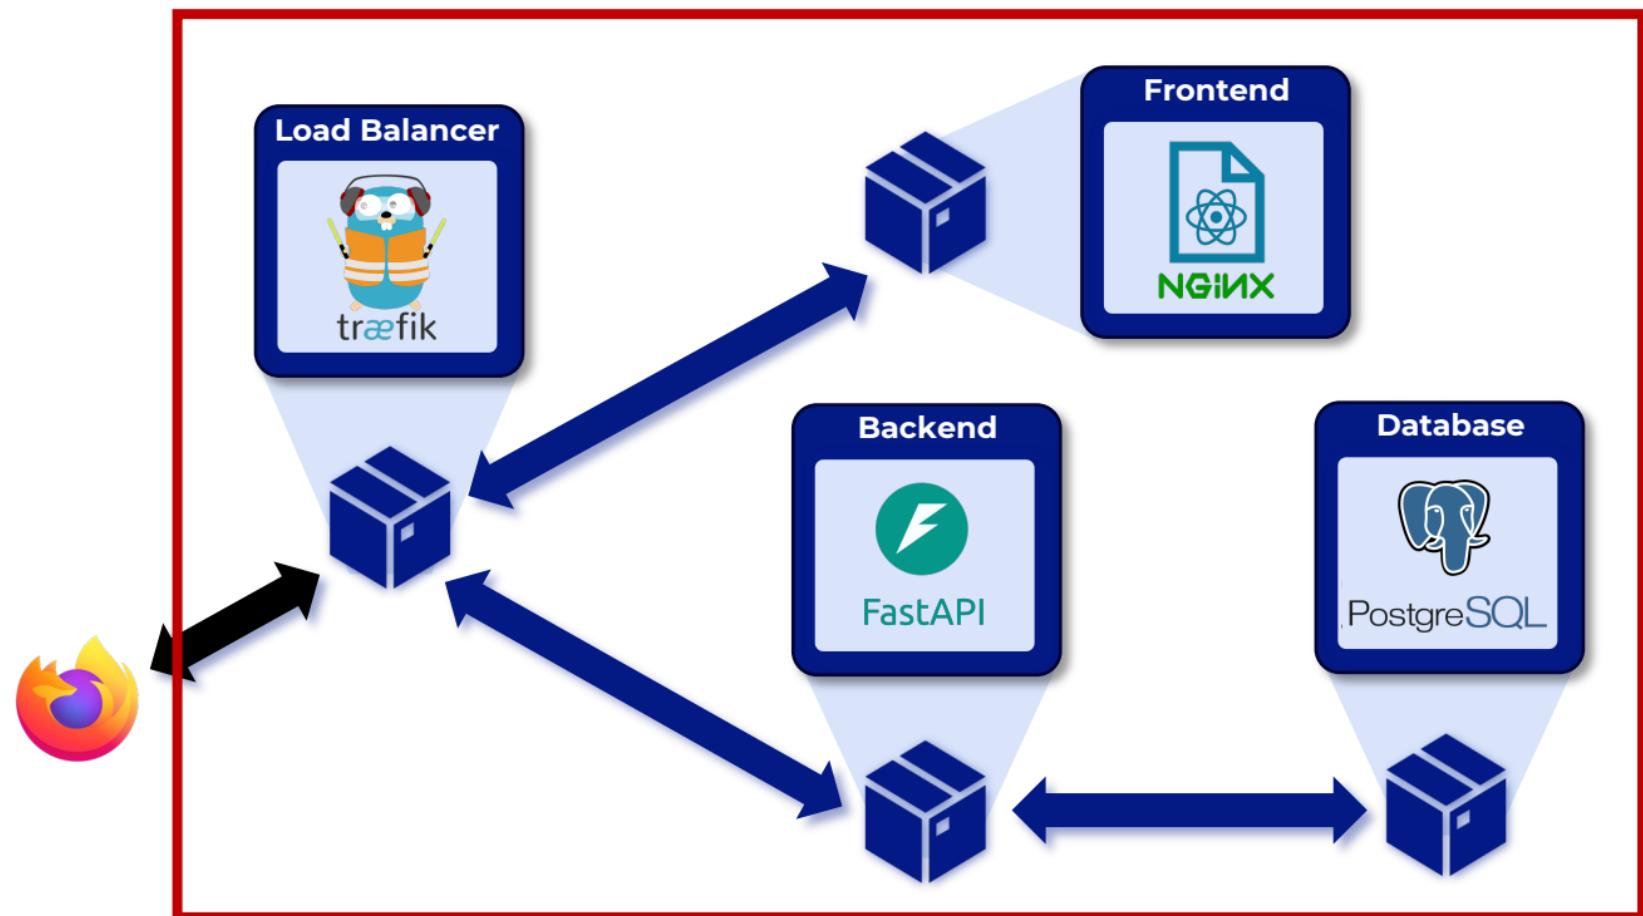
Application layer

Dependency layer

Base layer







Docker: The Efficient

Crash course: Linux shell



What do you have to know?



Navigate the filesystem

`ls`
`cd`
`pwd`



Manage files

`cat`
`touch`
`mv`
`cp`
`rm`



Inspect processes

`top`
`htop`



Use package managers

`apt`



Edit config files

`nano`



Network

`ping`
`curl`

Docker: The Efficient

How to install Docker?





Open-source

Docker Daemon:

Creates and runs containers

Docker CLI:

Provides the docker terminal command



Docker Inc.

Docker Inc. leads the development of the open-source docker components

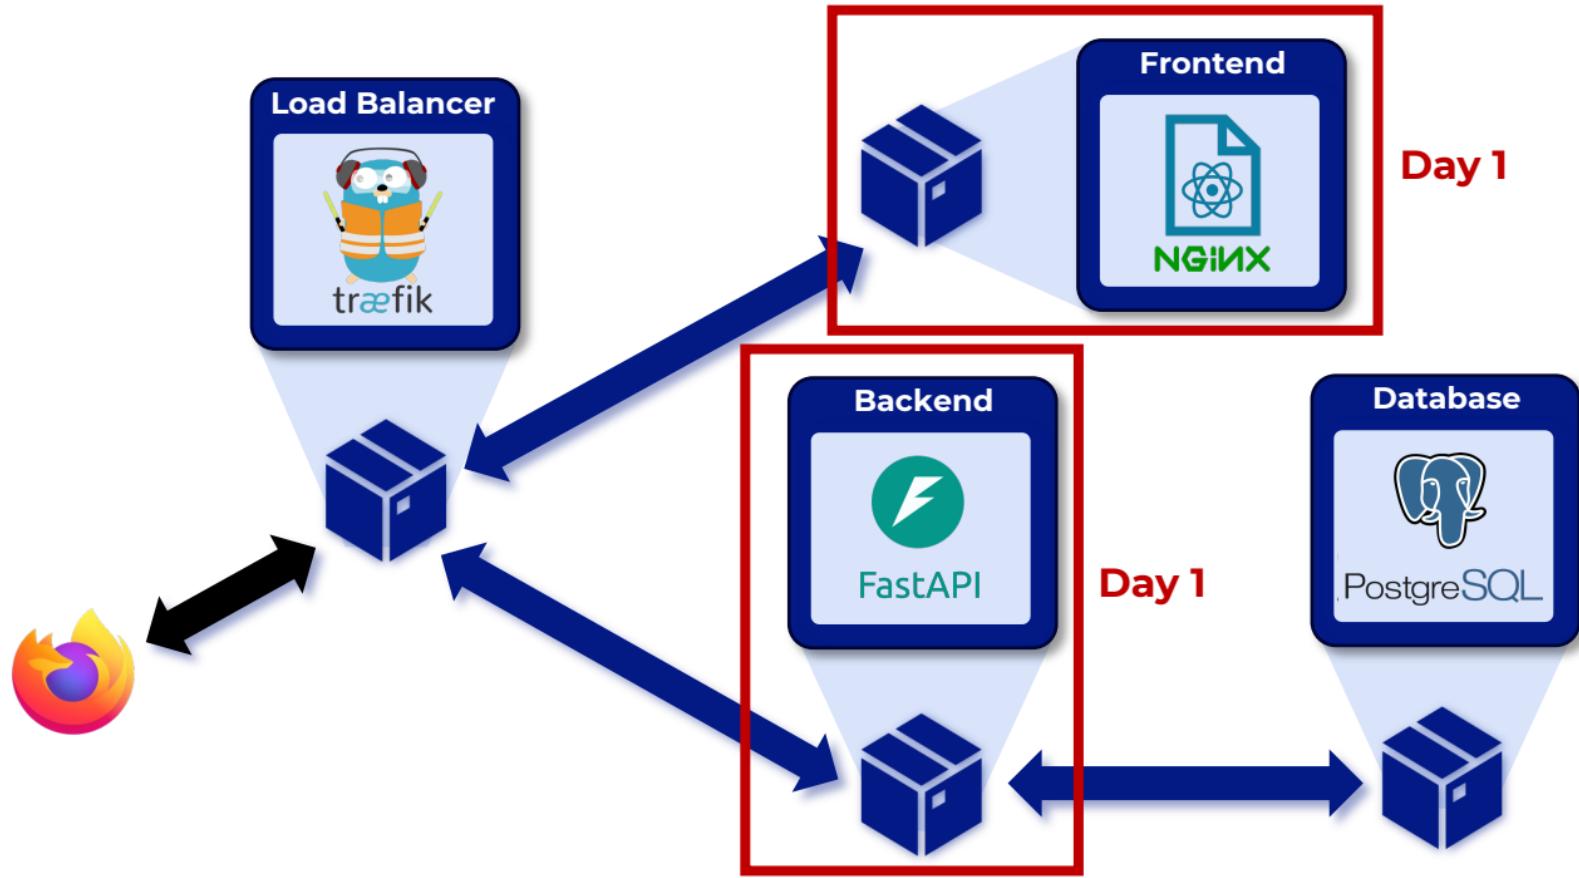
And offers an easy-to-use GUI (proprietary):

Docker Desktop

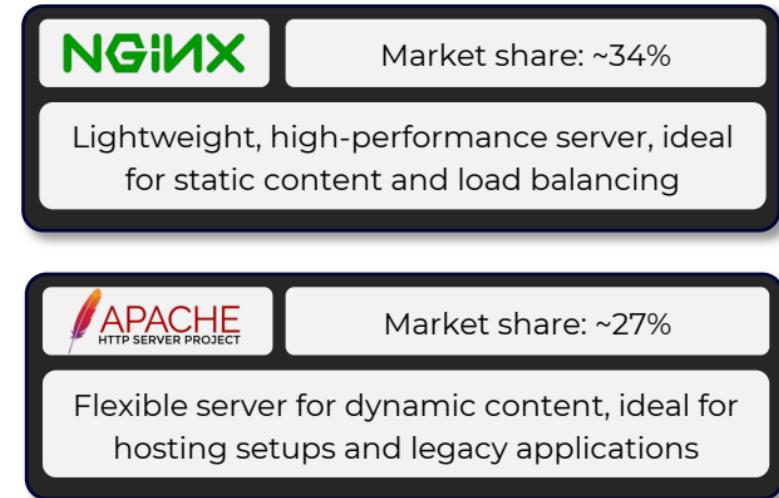
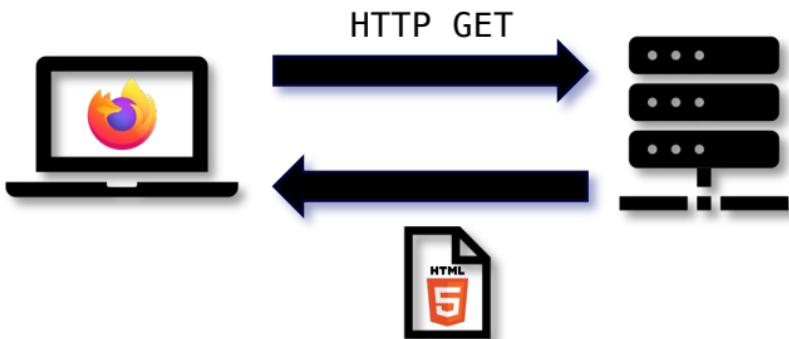
Docker: The Efficient

Making containers accessible (networking)





What are web servers?



IP addresses & port numbers



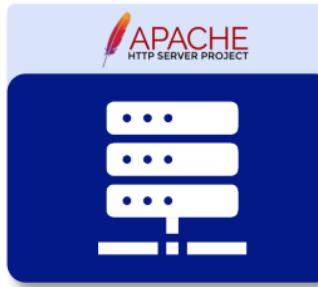
IP address

A unique identifier for devices in a network

Allows communication between devices



203.0.113.10



90.186.100.143

IP addresses & port numbers



IP address

A unique identifier for devices in a network

Allows communication between devices



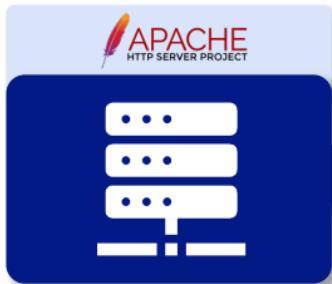
203.0.113.10 : 52345



Port number

Logical endpoint on a device

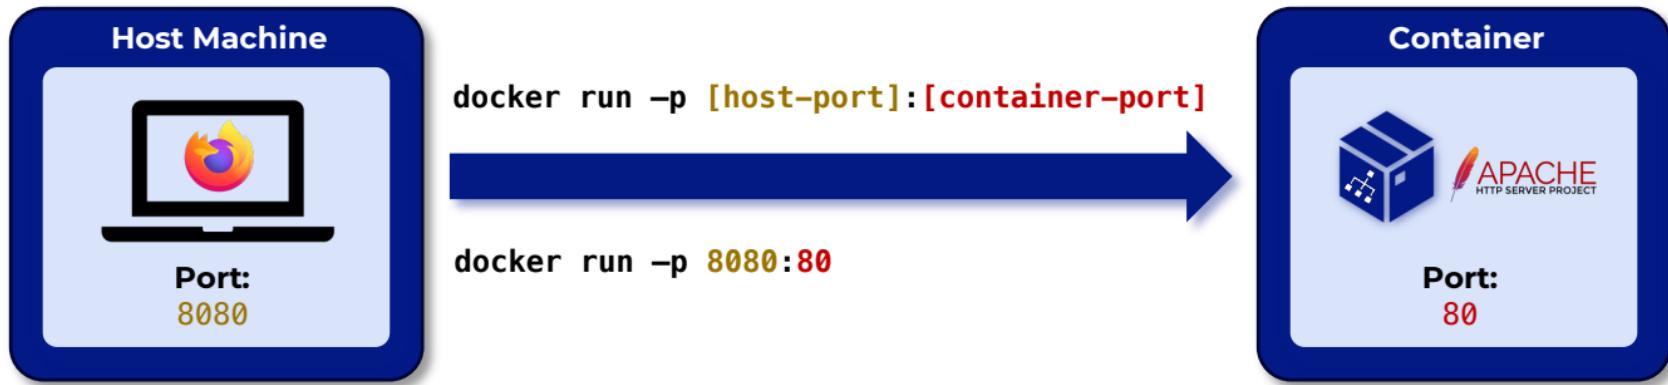
Ensures data is delivered to the correct application or service on that device



90.186.100.143 : 80



Port forwarding: -p



Docker: The Efficient

Running containers in the background: Interactive vs. detached





Interactive mode (-it)

Runs the container in the foreground and attaches a terminal (-t)

Allows direct interaction with the container (-i)



Detached mode (-d)

Runs the container in the background

This allows us to continue using our terminal



Access containers (docker exec)

We can use docker exec to execute commands

```
docker [container]  
exec -it  
<container>  
<command>
```



Challenge

Change the default HTML file in an NGINX container

Run an `nginx` container in the background, map its port to your local port `8000`

Make sure the port is not already in use; if it is, either free it or choose a different one.

Access the running container and start a `bash` shell

Edit the default HTML file `index.html` in `/usr/share/nginx/html` and change `<html>` to `<html style="background: plum">`

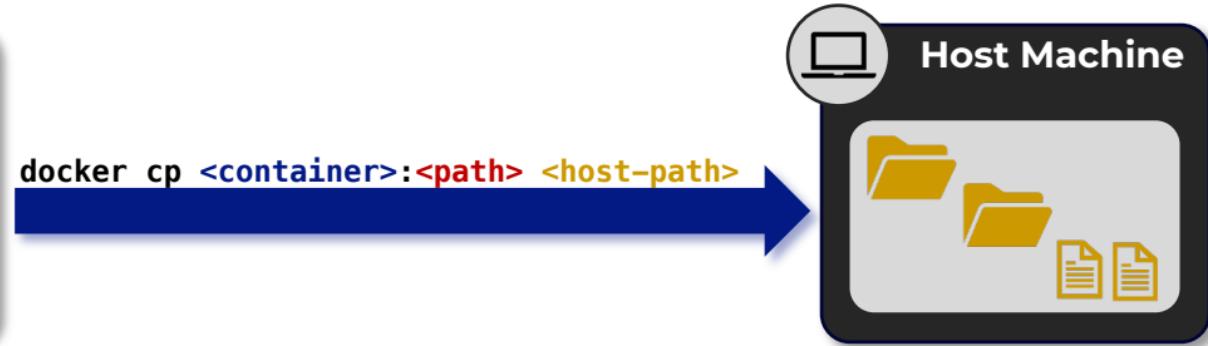
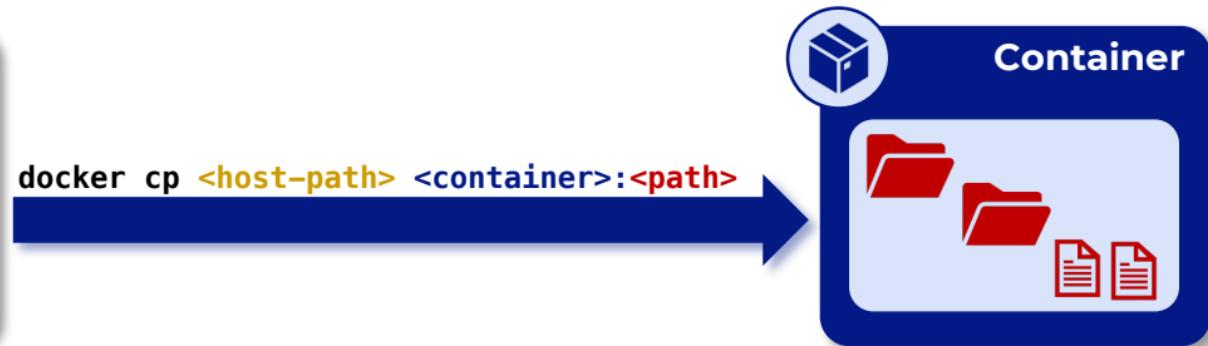
You may want to install a text editor like `nano` (through `apt`)

Verify the change in your browser – what has changed?

Docker: The Efficient

Copying data into / from containers





Docker: The Efficient

Challenge: Copying data into a container





Challenge

The goal: Deploy the expense tracker frontend with nginx

Start an nginx container

Copy the `/usr/share/nginx/html/index.html` file from the container to a local directory (as a backup)

Replace the contents of the `/usr/share/nginx/html` directory with the contents of the `dist` directory

Verify the change in your browser

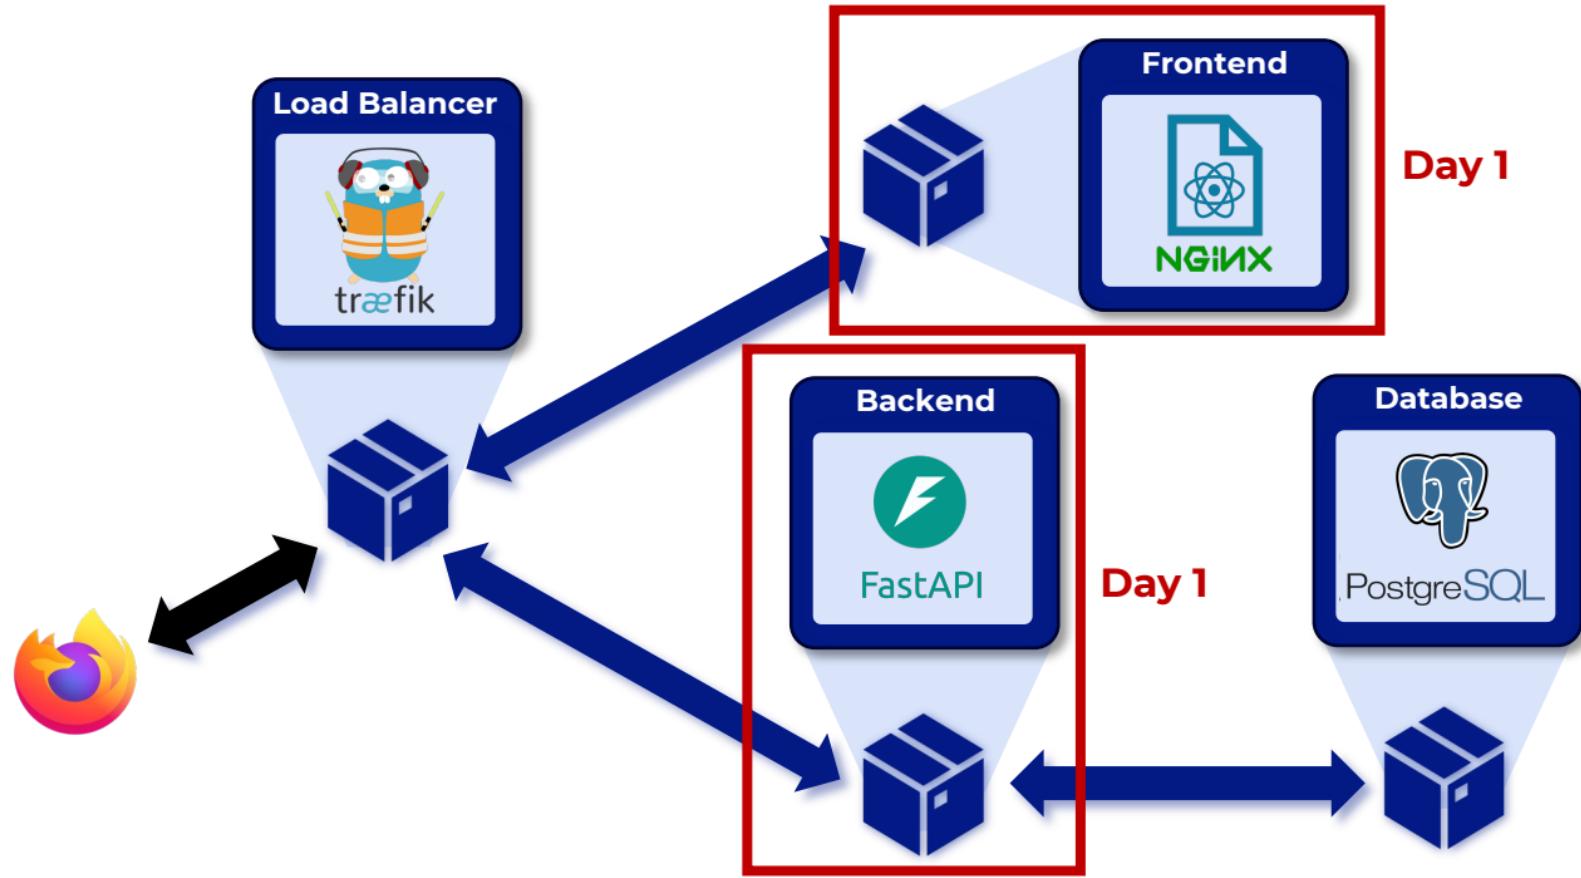
You can also specify a name, running mode, and port

The `dist` directory contains the production build of our react app. You can find it in the next lecture.

Docker: The Efficient

Overview: What are images?





Docker: The Efficient

What are images?



What is an Image?



Blueprint for containers

Provides everything needed to run an application

Includes code, runtime, libraries, environment variables, and other configurations



Immutable and read-only

The image itself cannot be changed

Changes occur only in the container's writable layer



Layered architecture

Images consist out of multiple layers to optimize storage

Layers are cached for faster builds and efficient reuse



Containers are the running instances created from images

Docker: The Efficient

Containers vs. images





Images

Serve as the template for creating containers



Provide read-only files that remain static and immutable



Define the filesystem, dependencies, and configuration to run a container
(but don't run on their own)



Persist and remain reusable across multiple containers and hosts



Containers

Act as active runtime instances created from an image

Add a writable layer for runtime changes

Execute the commands and processes defined in the image

Tie to the host they are created on, making them typically temporary

Docker: The Efficient

Finding images on Docker Hub



What is Docker Hub?



Centralized image registry

A platform to store, manage, and share Docker images

Hosts both public and private repositories



Collaboration and discovery

Enables teams to share images easily

Provides access to a vast library of official and community-contributed images



Integration with docker tools

Seamlessly integrates with the Docker CLI for pulling, pushing, and managing images

Acts as the default registry for Docker

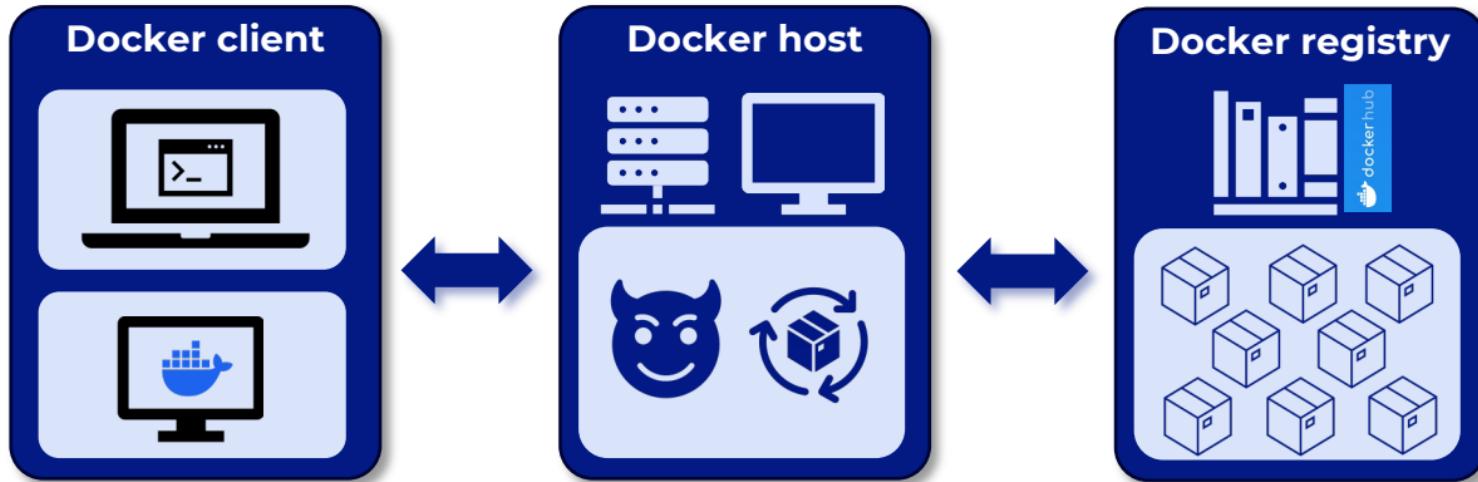


Docker Hub ensures consistency and accessibility for deployments

Docker: The Efficient

Dive deep into the container creation process!





Pull the image



`docker [image] pull`

Create a container



`docker [container] create`

Start the container



`docker [container] start`

`docker [container] run`

`-i`

`-t`

`-d`

`-p`

`-i`

Docker: The Efficient

Overview: Dockerfiles



Overview: Dockerfiles

► In this chapter:

- We'll explore Dockerfiles
- They allow you to create your own images
- ... which we can then use to execute our own application

► The structure:

► In the videos:

- We'll create a Dockerfile for the backend

► After that:

- You will create a first Dockerfile for the frontend (by using pre-compiled files)
- And we'll enhance this frontend together (and enabling compiling of .js files)

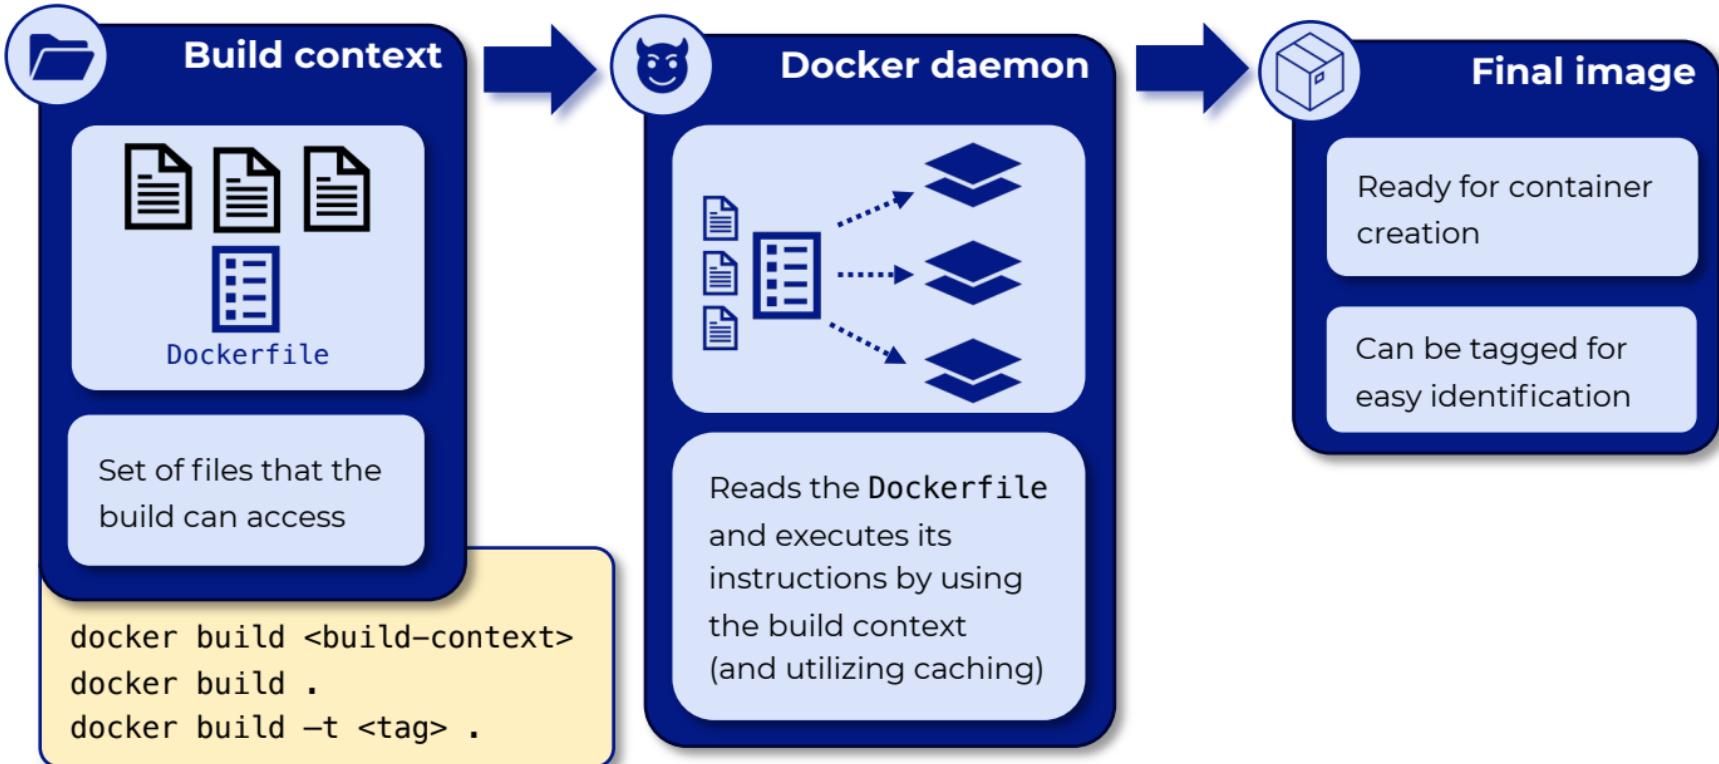
► We recommend:

- Using a text editor like Visual Studio Code:
- <https://code.visualstudio.com/>
- This will make it way easier to follow along

Docker: The Efficient

Build your first own image with a Dockerfile





The FROM Instruction



Definition

Specifies the base image for your custom Docker image

```
FROM <image>[:<tag>]
```

```
FROM python:3.12
```



Important notes

Defines the operating system or runtime environment for your app

Every Dockerfile must start with a FROM instruction

For now, a Dockerfile only uses a single FROM instruction



Best practices

Use official images from trusted sources

Choose lightweight base images for smaller builds when possible

Specify a tag to avoid unexpected changes

Docker: The Efficient

Choosing a suitable base image



How to choose a good base image?



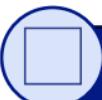
Runtime-specific

Includes runtime dependencies for faster app development



General-purpose & flexible

Ideal for apps needing a more specialized environment



Lightweight & minimal

Best for simple applications where size matters



Beware of Alpine



Advantages

A lightweight Linux distribution (~5 MB), that reduces image size

Minimal dependencies for faster builds



Limitations

Musl C Library

Smaller but not fully compatible with all software

Slower for apps with complex I/O or memory workloads (e.g., Python)

APK Package Manager

Limited Ecosystem

Less optimized packages



Best Practices

Ideal for lightweight apps like microservices or static binaries

Test thoroughly for third-party library compatibility

Consider `debian:slim` for better compatibility

Docker: The Efficient

The COPY & WORKDIR instructions



The COPY instruction



Definition

Copies files or directories from the host to the image during the build process

```
COPY <source> <destination>
```

```
COPY . /app/
```



Important notes

COPY only copies files that are part of the build context

Files from outside cannot be accessed



Best practices

Copy only necessary files to reduce image size and build time

Use a `.dockerignore` file to exclude unwanted files

Avoid copying sensitive files to improve security

The WORKDIR instruction



Definition

Sets the working directory within the container's filesystem for subsequent instructions



Important notes

WORKDIR persists in the container (and will be the default directory)



Best practices

Use WORKDIR to simplify commands and avoid repetitive paths

`WORKDIR <path>`

`WORKDIR /app`
`COPY . .`

Automatically creates the directory if it doesn't exist

Docker: The Efficient

The RUN instruction



The RUN Instruction

Definition	Important notes	Best practices
Executes a command during the image build process to install software, configure the environment, or prepare the image	Creates a new layer for each RUN instruction Layers are cached to speed up future builds if unchanged	Combine multiple commands into a single RUN instruction to minimize layers Try not to create files that are not needed – for example, by disabling caching
<code>RUN ["executable", "param1", "param2"]</code>	Exec format	
<code>RUN executable param1 param2</code>	Shell format	
		<p>Runs the command via <code>/bin/sh -c</code>, enabling shell features like pipes, environment variable expansion,...</p>

Docker: The Efficient

The CMD instruction



The CMD instruction



Definition

Specifies the default command that is executed when the container launches



Important notes

Later CMD instructions overwrite earlier ones



Best practices

Use the exec format whenever possible (it's not always)

`CMD ["executable", "param1", "param2"]`

Exec format

`CMD executable param1 param2`

Shell format

Runs the command via `/bin/sh -c`, enabling shell features like pipes, environment variable expansion,...

Docker: The Efficient

The EXPOSE Instruction



The EXPOSE instruction



Definition

Declares the port(s) that the container will listen on at runtime

`EXPOSE <port>[/protocol]`

`EXPOSE 80`

`EXPOSE 80/tcp`



Important Notes

Does **not publish the port** to the host automatically

To make it accessible on the host, use the `-p` or `--publish` flag with `docker run` the port



Best Practices

Use EXPOSE to document the container's intended ports for clarity and maintainability

Docker: The Efficient

Multi-stage builds



Problems with single-stage builds



Large image sizes

Build tools like `npm` or compilers remain in the final image

Unnecessary dependencies are deployed



Security risks

Build tools and extra layers expand the attack surface

Additional files may introduce vulnerabilities

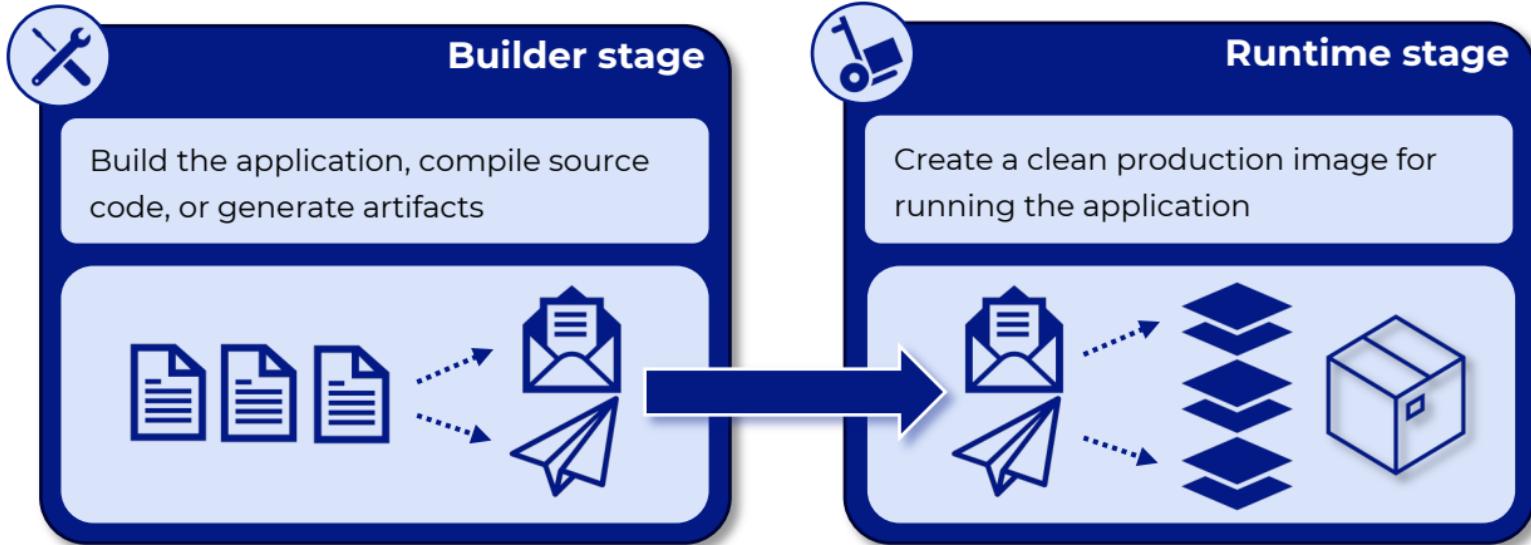


Inefficient builds

Minor source changes invalidate the cache, forcing full rebuilds

Managing separate `Dockerfiles` or scripts for building and serving increases complexity unnecessarily

Multi-stage builds



Docker: The Efficient

The ARG & ENV instructions



The ARG instruction



Definition

Defines build-time variables for the image build process



Important notes

Available during the build process only (not at runtime)

Pass values using the --build-arg flag



Best practices

Use ARG for build-specific configurations

Avoid using ARG for sensitive data

```
ARG <name> [=<default_value>]
```

The ENV instruction



Definition

Sets environment variables that persist across all subsequent instructions and into the running container



Important notes

Unlike ARG, ENV variables remain available at runtime after the container starts



Best practices

Use for runtime configuration

Avoid storing sensitive data directly in ENV

```
ENV <key>=<value>
```

Docker: The Efficient

The LABEL instruction



The LABEL instruction



Definition

Adds metadata to an image (key-value pairs)

`LABEL <key>=<value>`

`LABEL version="1.0.0"`



Important notes

An image can have more than one label

Labels included in base images are inherited by derived images

Enables searching and filtering images using `--filter label=<key>`



Best practices

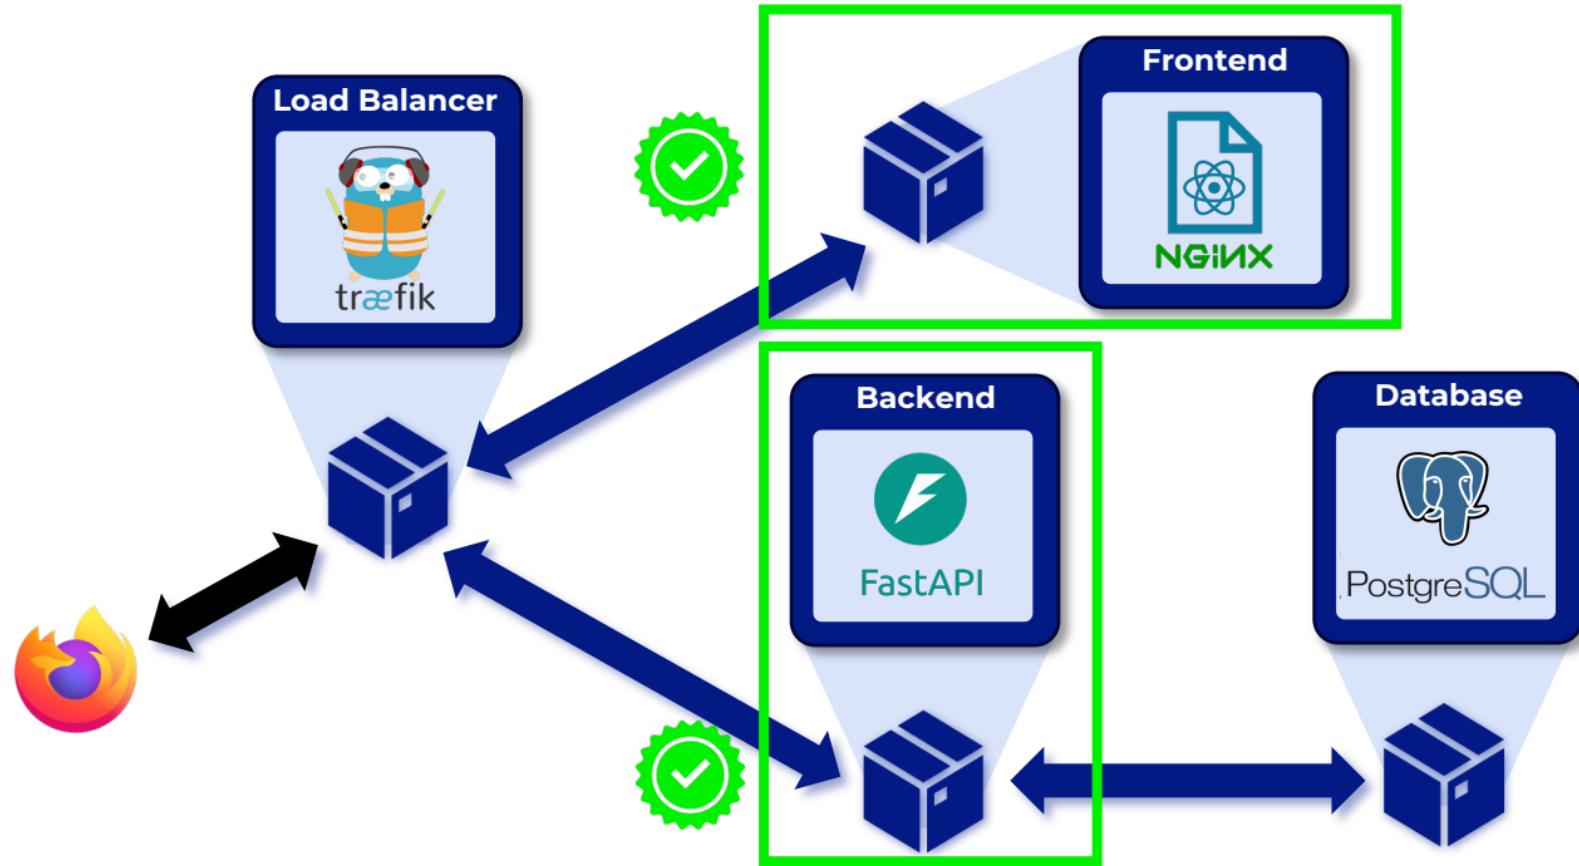
Follow consistent naming conventions

Keys should only contain lowercase alphanumeric characters and:
`. , _ , / , -`

Docker: The Efficient

Day 1: Conclusion

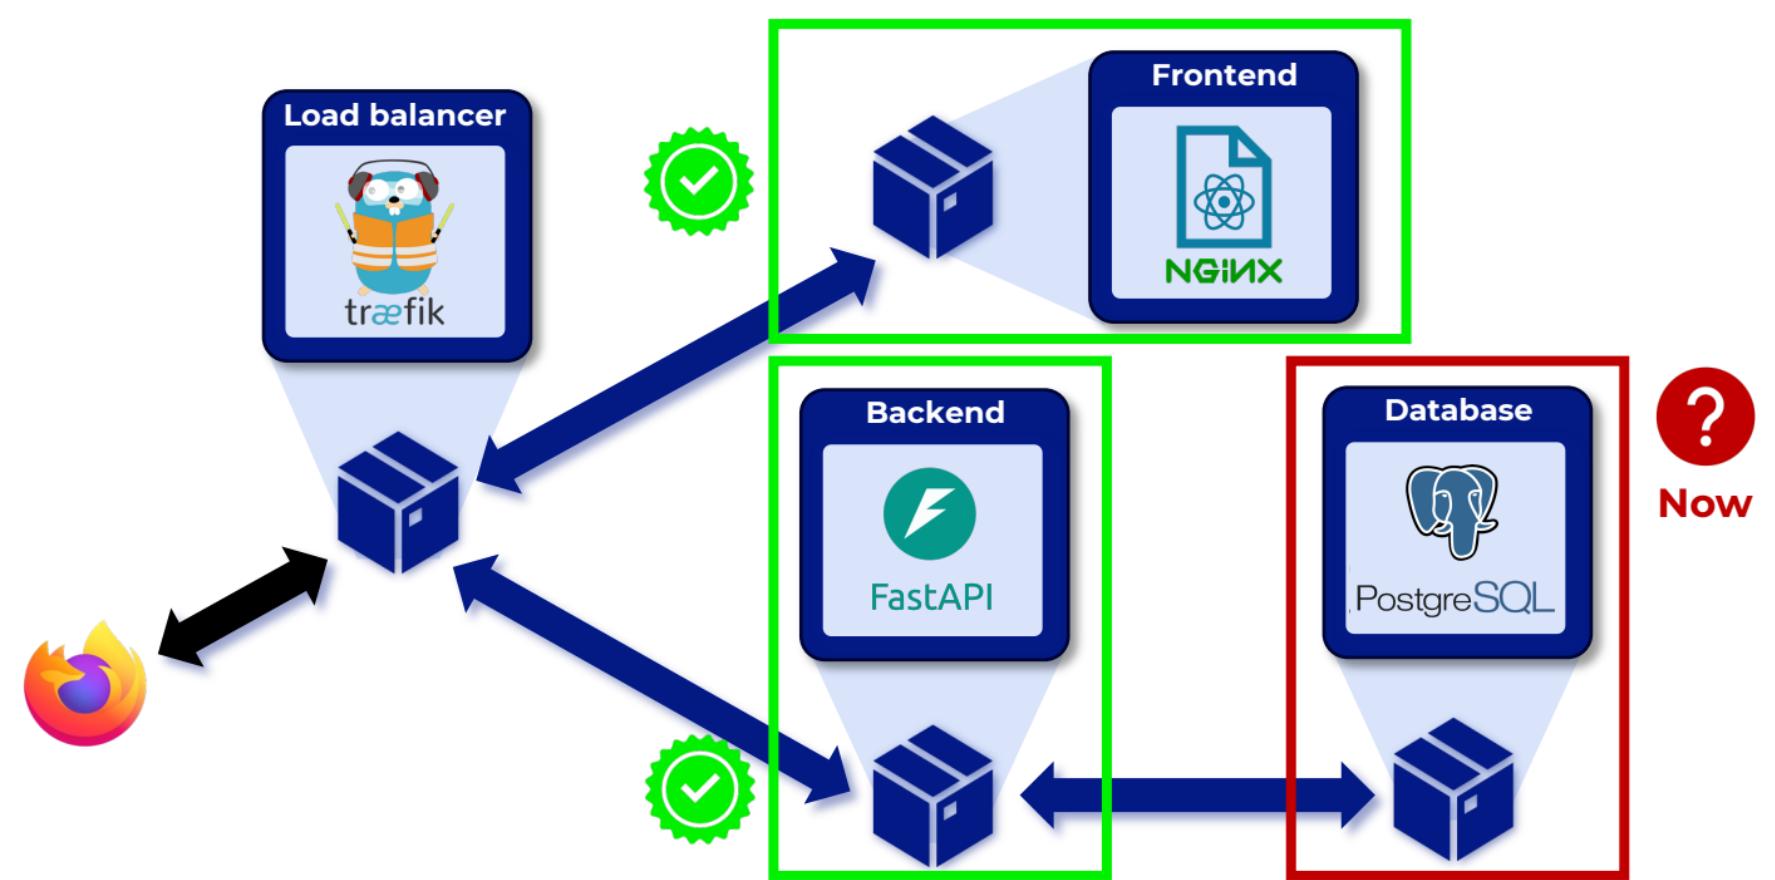




Docker: The Efficient

Persistent storage for docker containers







Data loss

Writable layers are ephemeral

When the container is removed: all stored data is lost



Limited scalability

Data inside the writable layer is tied to a single container

Cannot be shared between containers or accessed outside the container directly

Difficult to run multiple or new instances that share the same data

Docker: The Efficient

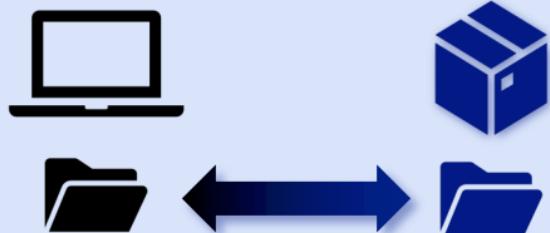
Storage mechanisms



Bind mount

Link a specific host directory or file to a container path

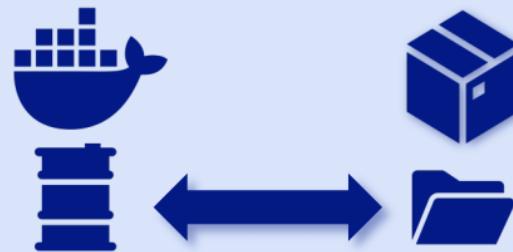
Directly reflects changes between host and container



Volume

Provide a dedicated storage space independent of the host's filesystem for persistent container data

Managed by Docker





Bind mounts

Tied to the host's filesystem, not managed by Docker



Depends on the host



Fully managed by Docker, independent of the host

Requires specifying the exact host path



Persistent; data remains even when container is removed

Limited portability; relies on host-specific path



Created and managed with Docker commands

Best for **development** (real-time file changes)



Works identically across any Docker-enabled system



Volumes

Ideal for **production** (faster, work identically everywhere)

Docker: The Efficient

Bind mounts





Definition

A mechanism to link a host directory or file to a container

Changes made on the host or in the container are reflected in real-time

```
docker run -v /host/path:/container/path[:access] <image>
```

```
docker run --mount  
type=bind,source=/host/path,target=/container/path[,access]  
<image>
```



Important notes

Docker itself does not offer any means to manage bind mounts

Files depend on the host's lifecycle



Best practices

During development, or for infrequently accessed files (config files)

Configure the container's access (e.g., read-only: ro)

Use absolute paths for the path on the host

Docker: The Efficient

Persistent storage through volumes





Definition

A Docker-specific resource to provide persistent storage beyond a container's lifecycle

Must be attached to a container during creation



Important notes

Some images automatically create volumes for you

You can manage volumes using `docker volume ...`



Best practices

Use in production for reliable and managed data storage

Explicitly name volumes for easier management

```
docker run -v <volume-name>:/container/path <image>
```

```
docker run --mount type=volume,source=<volume-name>,  
target=/container/path <image>
```

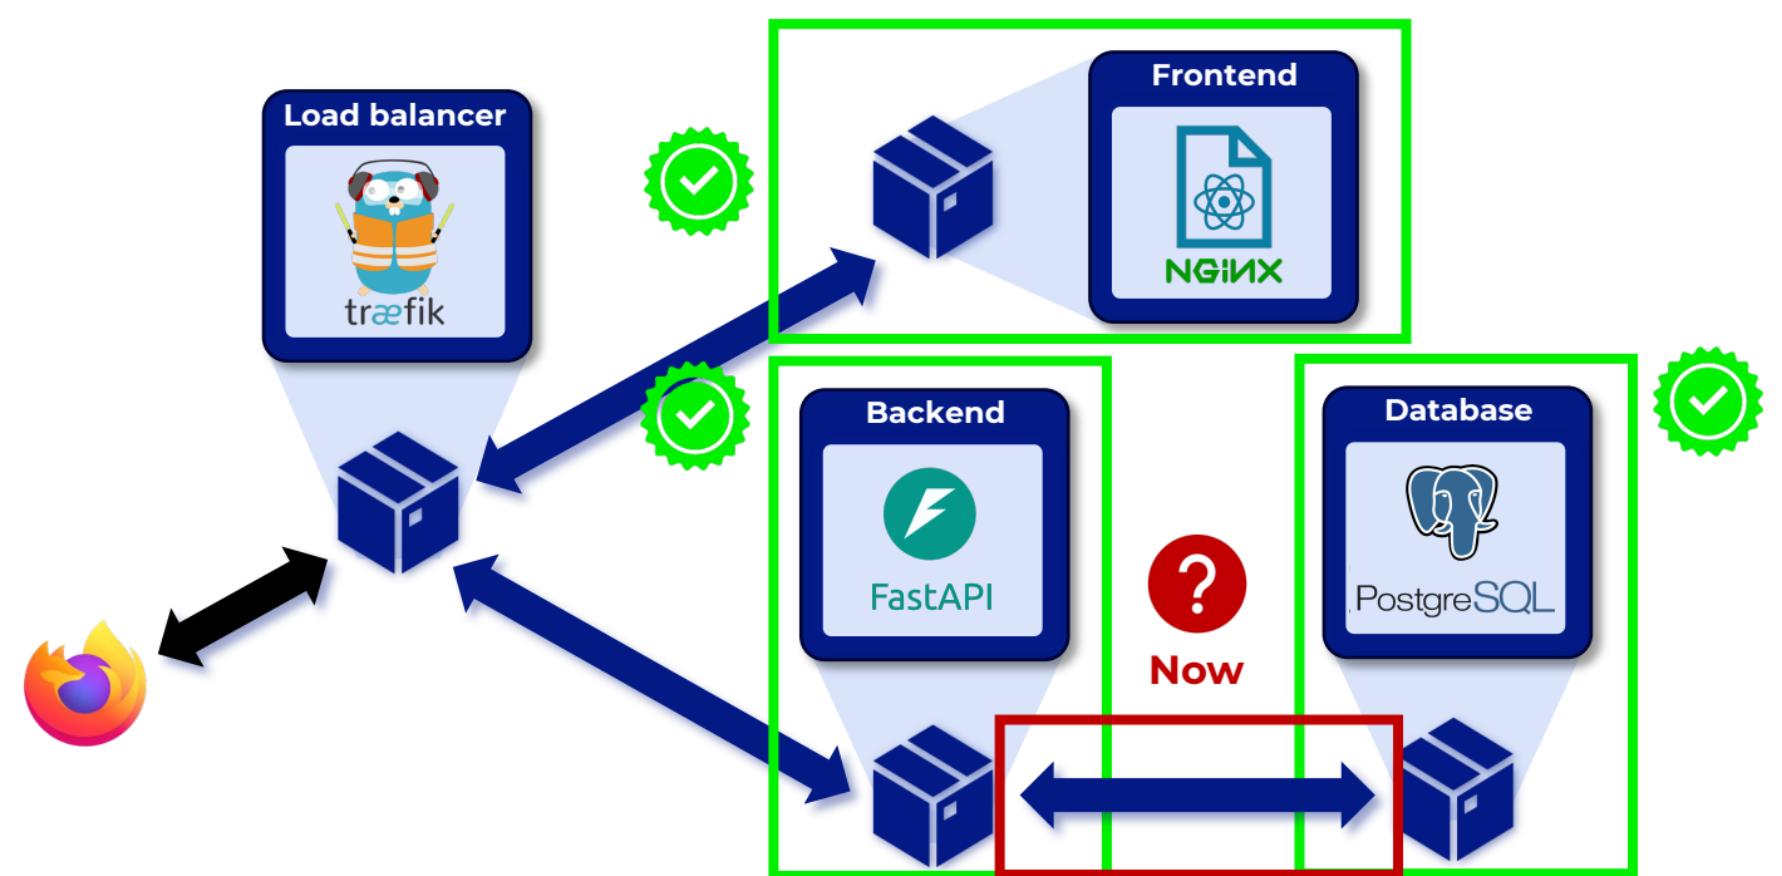


Docker: The Efficient

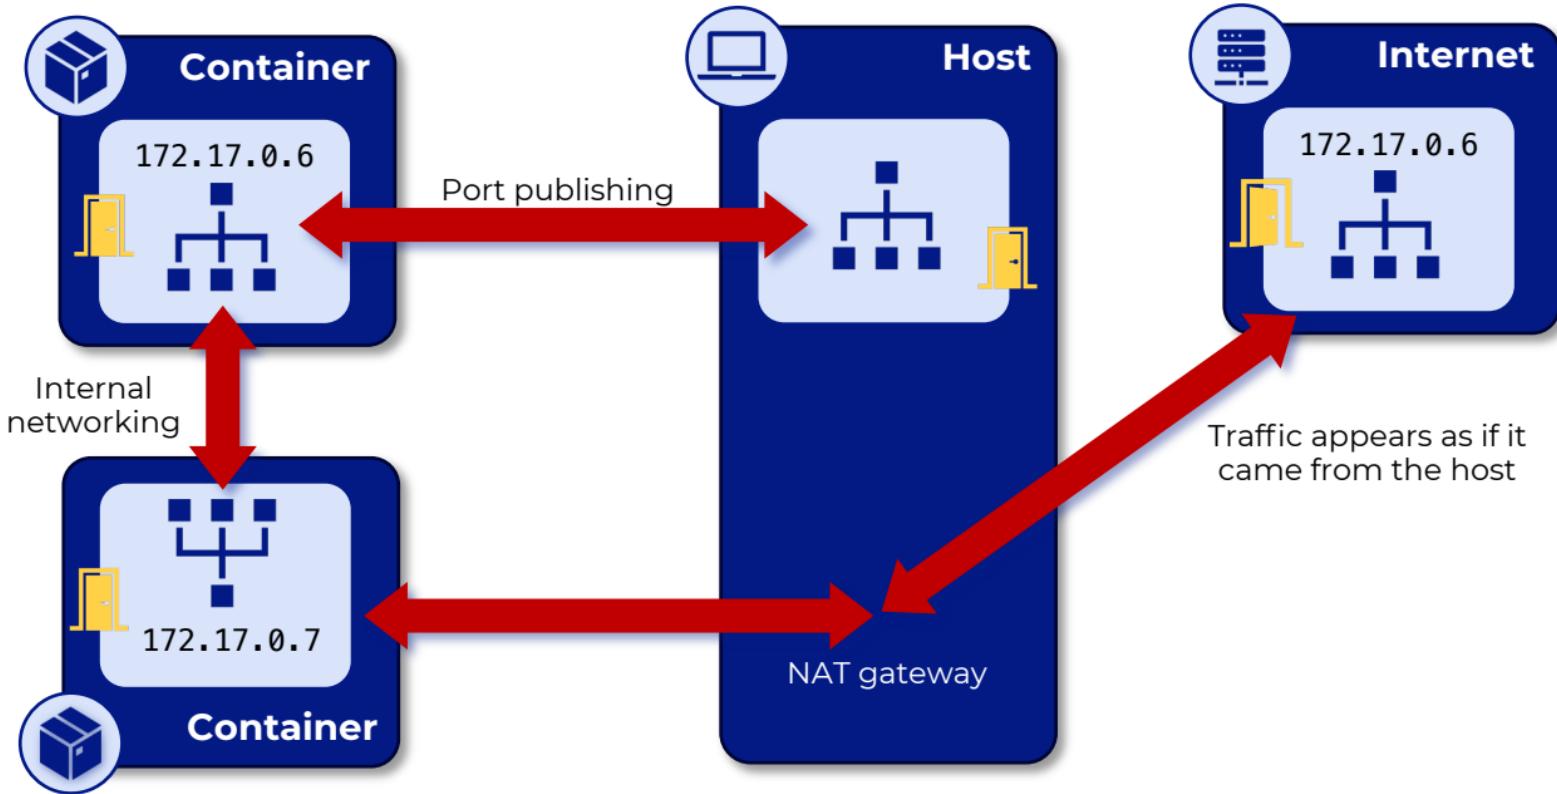
Container networking

[Day 2]





Container communication



Docker network modes



Bridge

Each container gets its own isolated network stack, including an IP address

Containers can communicate with each other using their IP

Default mode



Host

Removes any network isolation from the host

The container shares the host's IP address and uses its ports directly

Leads to low latency

Only available on Linux



None

The container has no network connectivity or IP address

Completely isolated from any network

Use for highly secure setups or manual configuration

Docker: The Efficient

Creating a custom network



Custom networks



Definition

You must create a network before assigning containers



Important notes

Containers on the same **custom** network can communicate via their container names



Best practices

Use custom networks instead of relying on the default bridge for multi-container setups

```
docker network create <name>
```

```
docker network connect <network-name> <container-id>
```

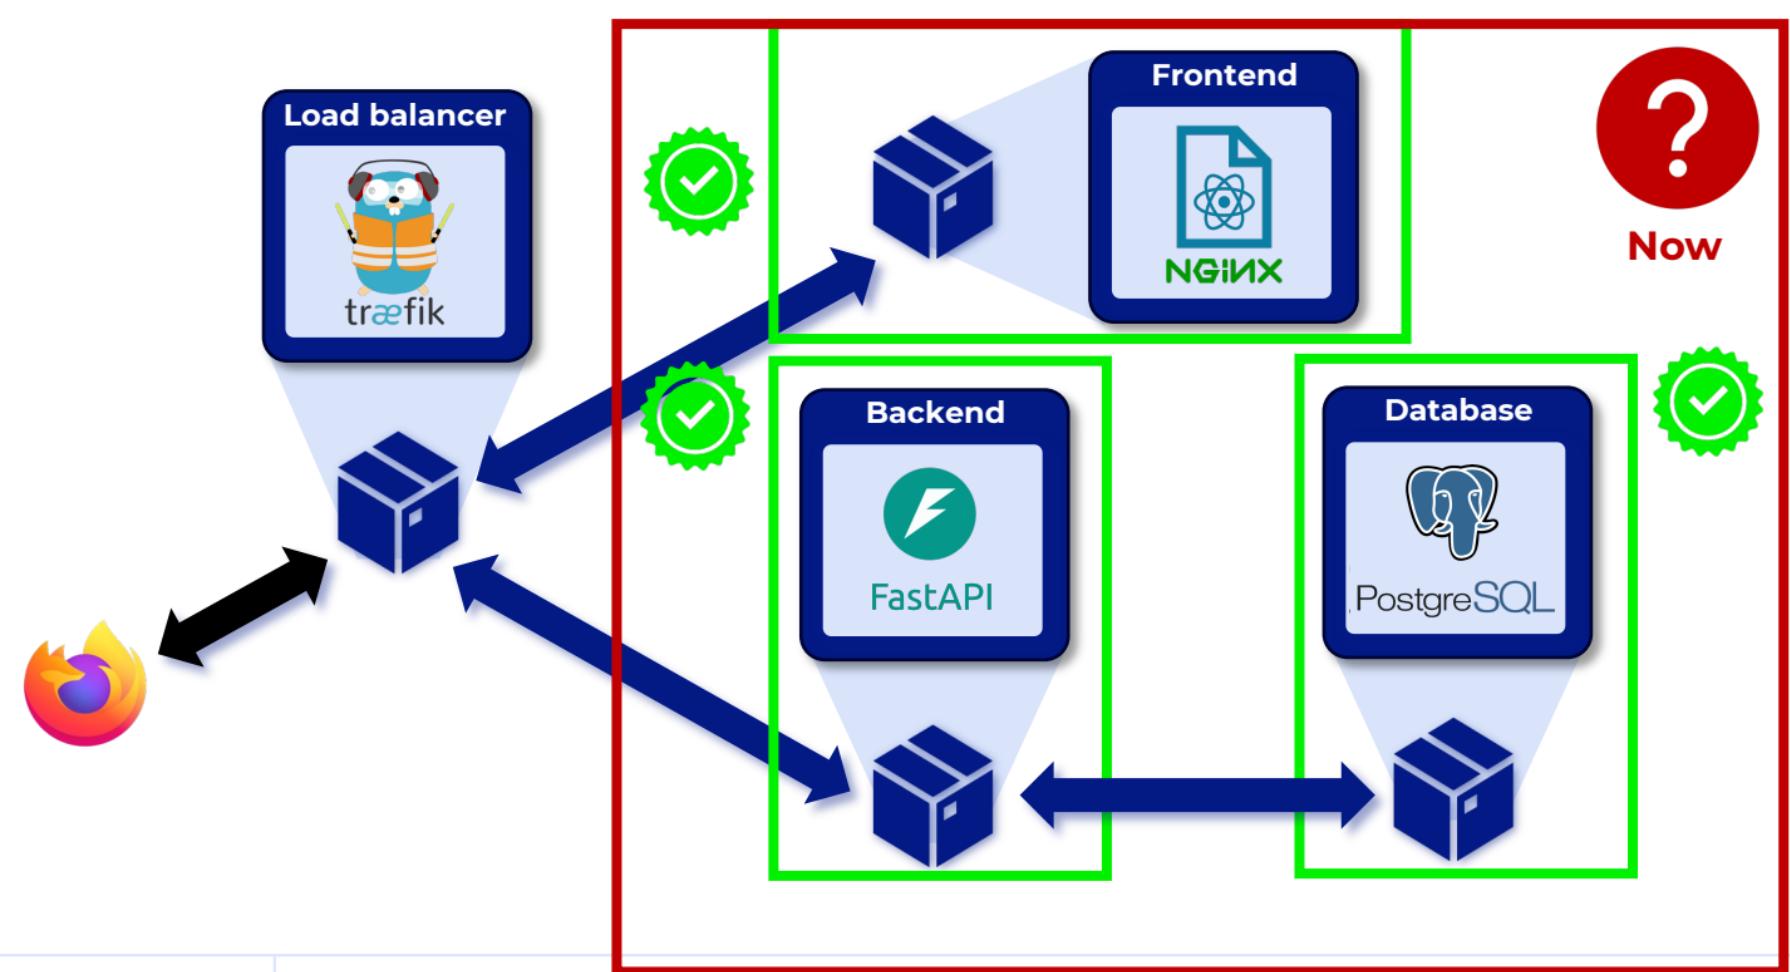
```
docker run --network <name> <image>
```

Create separate networks per stack to avoid naming conflicts and improve security

Docker: The Efficient

Launching the application





Docker: The Efficient

The tool docker compose

[Day 2]



Manual setup problems



Configuration complexity

Each container must be started individually with long docker run commands

Requires manually setting environment variables, volumes, and networks



Reproducibility issues

Difficult to ensure the same configuration across multiple machines

Requires extensive documentation for setup and dependencies

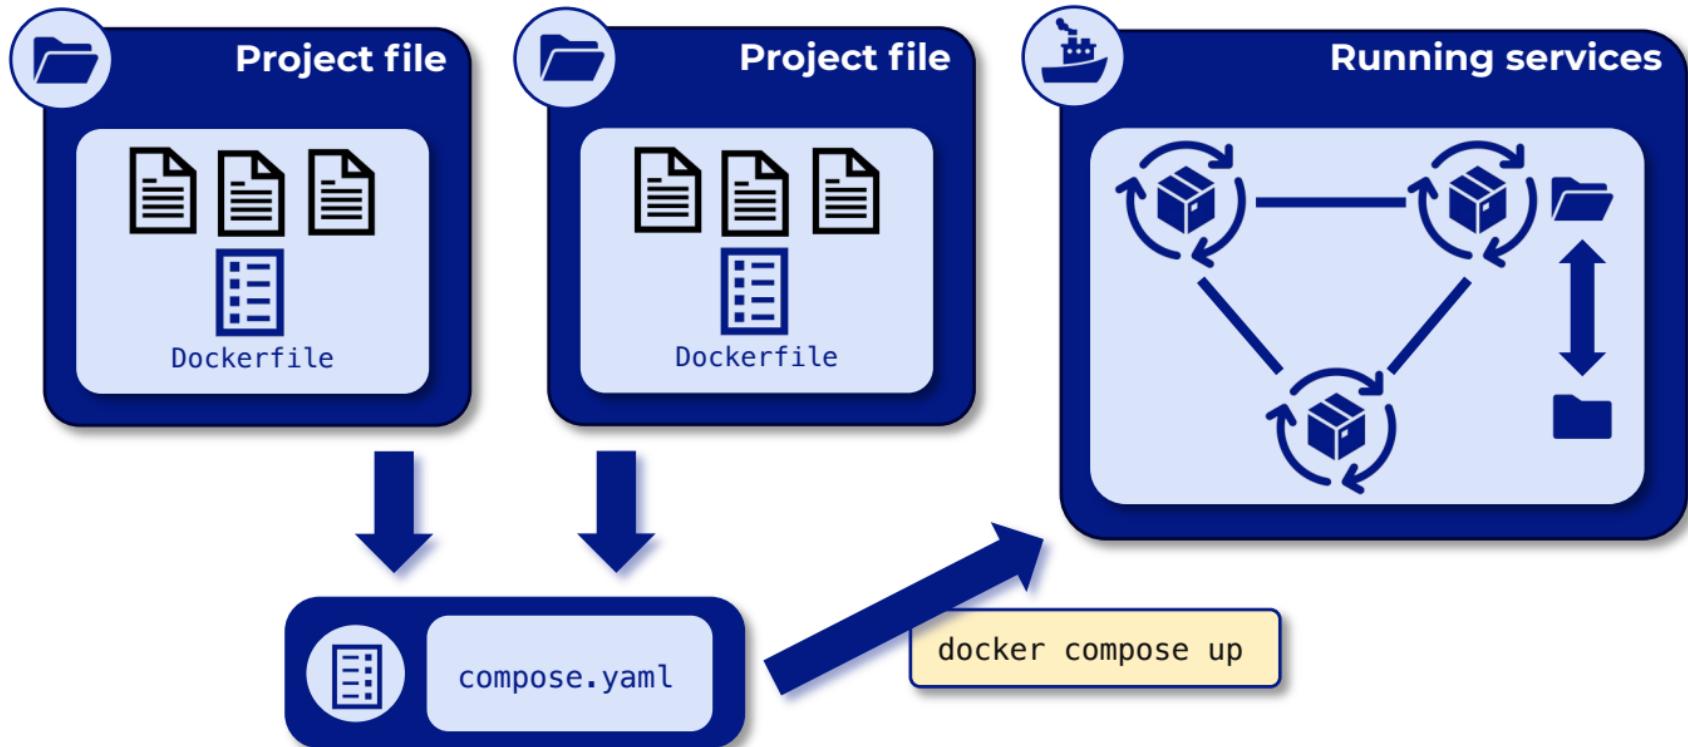


Scalability limitations

Scaling services requires manually running and managing multiple containers

No built-in way to define and maintain multi-container relationships

How docker compose works



Docker: The Efficient



Writing a `compose.yaml` file



YAML Format

Human-readable configuration format

Relies on indentation and key-value pairs



Key Sections

services:

Specifies functional units, grouping containers by their roles (e.g., app, db)

networks:

Defines custom networks, lists all networks used in services

volumes:

Defines persistent storage, lists all named volumes used in services



Benefits

Centralizes configuration for multi-container setups

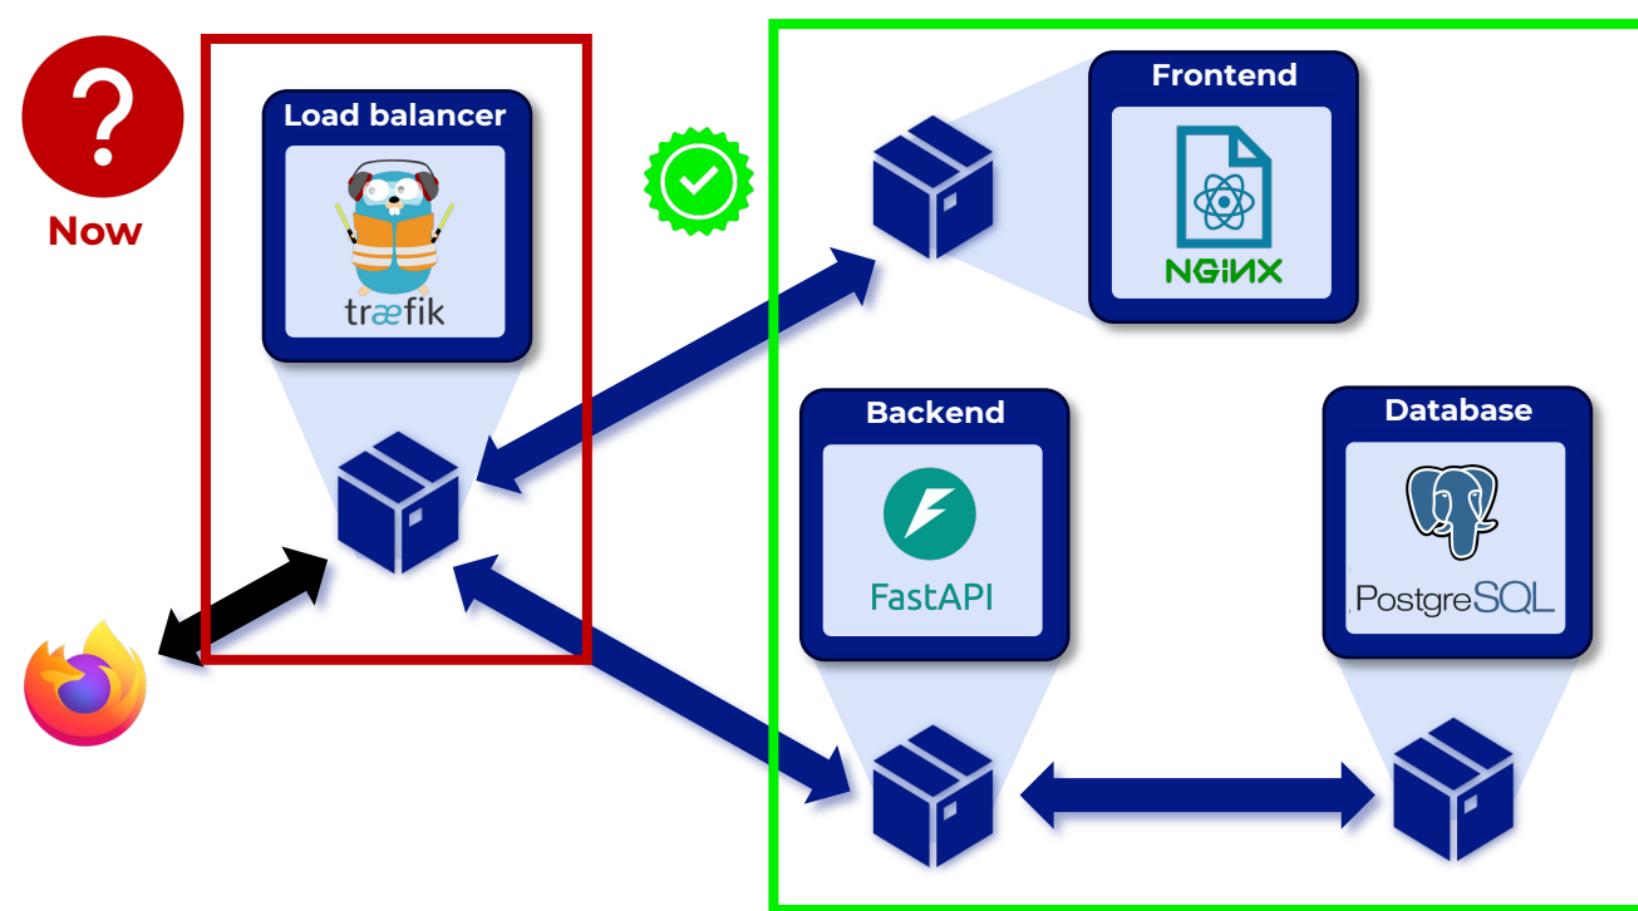
Ensures reproducibility across different environments

Simplifies deployment with a single command

Docker: The Efficient

Traefik: Load balancer





Docker: The Efficient



Day 3: Using docker in daily workflows

Overview

- ▶ Create a combined development / deployment setup in docker compose
- ▶ **Private registries & repositories:**
 - ▶ Where can we store our images?
 - ▶ Can we configure an automated deployment setup?
 - ▶ GitHub -> private registry -> server
- ▶ **How to deploy to a real server:**
 - ▶ How to connect to a remote server (ssh)
 - ▶ How to enable https?

Docker: The Efficient

Overview: How to distribute docker images

Overview

► **You will learn:**

- ▶ How to store an image in a private docker repository
- ▶ This is fundamental when deploying to production
- ▶ You will learn how to handle multi-architecture builds
(arm64, amd64,...)

► **This will lay the groundwork for the next section:**

- ▶ Automating the entire deployment workflow

Docker: The Efficient

Storing an image in a private repository





Repositories

A collection of related images



Stores multiple versions (tags) of an image



Example: nginx, python **or** php



Can be public or private



A service that stores and manages repositories

Hosts multiple repositories for different images

Docker Hub, GitHub Container Registry, AWS ECR

Can be cloud-based or self-hosted



Registries